

# Computer Systems Principles

## C Structures



# Learning Objectives

- To learn and apply C structures
- To understand a little about alignment
- To learn and apply C enums
- To learn and apply C unions
- To understand and apply C typedef

# C Structures

- **Essential**
  - For building up interesting data structures
- **Definition**
  - A **C structure** is a collection of one or more variables, typically of different types, grouped together under a single name for convenient handling
  - Kind of like a Java class with public instance variables and no methods

# C struct

- **Defines a new type**
  - A new kind of data type that the compiler regards as a unit or aggregate of variables/types.
- **Example:**

```
struct Date {  
    int day;  
    int month;  
    int year;  
};
```

# Structure Properties

- Individual components of a struct type are called **members** (or **fields**).
- Members can be of **different types** (primitive, array, or struct)
- A struct is *named* as a whole while individual members are named using field identifiers
- Complex data structures can be formed by defining *arrays of structs*.

# More struct Examples

- **Examples:**

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentGrade {  
    char name[25];  
    char course[9];  
    int lab[5];  
    int homework[7];  
    int exams[2];  
};
```


# Declaring a struct Variable

- Declaration of a variable of struct type:

**<struct type> <identifier list>;**

- Example:**

```
struct StudentRecord student1;
```



# student-01.c example

- **Let us compile this example**
  - What do you noticed about the output of this program as compared with the previous slide?
  - Is the size of this struct the same as we predicted?
  - Why is this or is this not the case?



# Data Allocation and Alignment


- **Data Allocation**

- Each variable definition is allocated bytes in memory according the type of that variable
- e.g., char = 1 byte, int = 4 bytes, double = 8 bytes
- This is allocated in a special place in memory known as the **stack**.


- **Data Alignment**

- Machines are more efficient if allocated data is on a **word boundary**.
- A word is typically 4 bytes

# Alignment in Memory




# Alignment in Memory



Each box is a byte  
**and** has a location.

Memory is very much like a a  
giant character array!


# Alignment in Memory



So, how do we allocate the structure definition below?

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```


# Alignment in Memory



We allocate 25 bytes for the character array **name**. In this example we start from 0, however, we could start from anywhere in memory.

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```

# Alignment in Memory




But, machines are typically **more efficient** if data is allocated on the start of 4 byte “word boundaries” (every 4<sup>th</sup> byte) e.g., 4, 8, 12, 16, 20, 24, ...

That is, the starting memory index for allocating the next data type should be  $(\text{index} \% 4) = 0$

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```


# Alignment in Memory



So, the compiler will make the allocation of your structures more efficient by **padding** the bytes so the **next allocation** will be **4-byte aligned**!

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```

# Alignment in Memory



So, the compiler will make the allocation of your structures more efficient by **padding** the bytes so the **next allocation** will be **4-byte aligned**!

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```




# Alignment in Memory



Next, we allocate bytes for the next type – this is a 4-byte integer, so it is already aligned properly.

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```


# Alignment in Memory



Next is the character. Will the next allocation be aligned properly?

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```


# Alignment in Memory



Nope! So the compiler will pad this out with 3 additional bytes

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```

# Alignment in Memory



Lastly, we allocate the last field in the structure which is an 8-byte double which is 4-byte aligned.

So, instead of 38 bytes – we get 44 bytes!

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1;
```

# What if...

**We change around the fields in a structure?**

Let us take a look at this example:

student-01-alt.c

# iClicker Question

The size of this struct will be:

- A. 38 bytes
- B. 40 bytes
- C. 44 bytes
- D. 48 bytes
- E. None of the above

# What Happened?

- Previously, we had 44 bytes.
- Now, we have 40 bytes allocated!
- Same fields ?!

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord {  
    int id;  
    double gpa;  
    char name[25];  
    char gender;  
};
```

Think about this for a moment and write down an explanation.  
I will randomly ask for four responses and read them!

# Rounding up ...

- The new layout uses 38 bytes, packing better than the old one, but ...
- The compiler rounds struct sizes up to a multiple of 4 bytes, giving 40 as the length.
- Why?
  - The main reason is that each struct in an array of these structs needs to *start* on a 4 byte boundary, so the individual structs need to be a multiple of 4 bytes in size.



# Structure Initialization

- **There are three ways to initialize a struct**
  - Positional initialization
  - Named initialization
  - Copy initialization
  - Initialize individual fields

# Positional Initialization

**Positional initialization** allows you to provide the values for each of the fields based on the position of each structure member:

```
struct StudentRecord student1 = {  
    "John Doe", 1234567, "M", 3.95  
};
```

# student-02.c example

- **Let us compile this example**
  - Notice that each value in the structure initializer is exactly the same position as the structure definition.
  - Look at how we use the ‘.’ operator to access the individual fields in a structure. This should be reminiscent of how you access fields in Java.

# Named Initialization

**Named initialization** allows you to provide the values for each of the fields based on the name of each structure member:

```
struct StudentRecord student1 = {  
    .id    = 1234567,  
    .gpa    = 3.95,  
    .gender = 'M',  
    .name = "Harry Potter"  
};
```

# student-03.c example

- **Let us compile this example**
  - Notice that each value in the structure initializer can occur in any position.
  - Do not forget to use the ‘.’ prefix for each field name in the initializer!

# Copy Initialization

**Copy initialization** allows you to initialize a structure by assigning an existing structure:

```
struct StudentRecord student1 = {  
    .id      = 1234567,  
    .gpa     = 3.95,  
    .gender  = 'M',  
    .name    = "Harry Potter"  
};
```

```
struct StudentRecord student2 = student1;
```

# student-04.c example

- **Let us compile this example**
  - Notice that we can initialize by using the assignment operator '='.
  - This will automatically copy the memory from the structure on the right-hand side to the structure on the left-hand side.

# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord student1;
```

```
student1.id          = 1234567;
```

```
student1.gender      = 'M' ;
```

```
student1.gpa         = 3.95;
```



# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord student1;  
  
student1.id          = 1234567;  
student1.gender     = 'M' ;  
student1.gpa        = 3.95;  
student1.name       = "Harry Potter";
```

**What about this one?**

# student-05.c example

- **Let us compile this example**
  - What problems do we encounter with this example?
  - Why can't we assign a string to a character array?
    - The string is a `char *` type
    - Arrays are not modifiable values, that is, you can't reassign them to "point" to different locations in memory.
    - Huh? This will be more clear when we talk about pointers 😊

# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord student1;  
  
student1.id          = 1234567;  
student1.gender      = 'M' ;  
student1.gpa         = 3.95;  
student1.name        = "Harry Potter";
```

**So, how do we fix this?**

# strncpy

- **Copying Strings**

- #include <string.h>

- A library for manipulating C strings

- To assign a new string value to a C string (e.g., character array) you must use the *strncpy* function to **copy** the bytes into the array.

# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord student1;
```

```
student1.id          = 1234567; Size of the destination
```

```
student1.gender      = 'M' ;
```

```
student1.gpa         = 3.95;
```

```
strncpy(student1.name, "Harry Potter", 25);
```



**We use the strncpy function!**

# student-05-fix.c example

- **Let us compile this example**
  - Notice that need to specify the number of bytes.
- **Why do I need to specify the number of bytes?**
  - There is also a *strcpy* function that does not require the number of bytes.
  - but, this function is dangerous because it is possible to copy a larger string into a smaller destination array **overwriting** adjacent memory!
  - This is called **buffer overflow** and can be used to exploit system vulnerabilities!



# student-06.c example

- **One last example...**
  - Creating a “constructor” function for creating new structures is a useful pattern in C.
- **Let us compile this example**
  - Notice that need to specify the number of bytes.
- **Some new things:**
  - `strlen` for computing the length of a string
  - Functions in C are “call by value”!
    - But, what about passing arrays to functions?

# Activity!

- **strlen(char s[])**  
**strncpy(char dest[], char src[], int n)**
  - Take a moment to implement these functions!
  - Work with the people around you!
  - Write it down on a piece of paper!

